

Week 14 IP

Ted Askoye

9/23/2020

1. Business Understanding

1 a.) Defining the Question

You are a Data analyst at Carrefour Kenya and are currently undertaking a project that will inform the marketing department on the most relevant marketing strategies that will result in the highest no. of sales (total price including tax). Hence making an analysis to Customer Data from a the supermarket and implement dimensionality reduction.

2. Defining the Metrics of Success

The success of this analysis will occur when we analysis the customer data to understand it fully and later implementing the appropriate dimensionality reduction techniques.

3. Context

Dimensionality reduction is the process of reducing the number of random variables under review, by getting a set of principal variables. It can be divided into feature selection and feature extraction and is important for the visualization of features while it also helps deal with multicollinearity of the features.

4. Experimental Design

We will define the question, the metric of success, context and experimental design taken. This will be followed by reading and exploring the dataset and its appropriateness of the available data to answer the given question. This will be followed by cleaning the data off outliers, anomalies and null values from missing data, perform an exploratory data analysis after which we will Implement feature extraction and feature selection, record our observations and provide a conclusion and recommendation.

5. Data Relevance

Our data is very relevant to our research question.

6. Loading relevant Libraries and Reading the Data

```
# Importing the required packages
```

```
library("data.table")
library("plyr")
library("dplyr")
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:plyr':
```

```
##
```

```
##      arrange, count, desc, failwith, id, mutate, rename, summarise,
##      summarize
```

```
## The following objects are masked from 'package:data.table':
```

```
##
```

```
##      between, first, last
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      intersect, setdiff, setequal, union
```

```
library("tidyverse")
```

```
## -- Attaching packages ----- tidyverse.
```

```
## v ggplot2 3.3.2      v purrr 0.3.4
```

```
## v tibble 3.0.3       v stringr 1.4.0
```

```
## v tidyr 1.1.2        v forcats 0.5.0
```

```
## v readr 1.3.1
```

```
## -- Conflicts ----- tidyverse.
```

```
## x dplyr::arrange() masks plyr::arrange()
```

```
## x dplyr::between() masks data.table::between()
```

```
## x purrr::compact() masks plyr::compact()
```

```
## x dplyr::count() masks plyr::count()
```

```
## x dplyr::failwith() masks plyr::failwith()
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::first() masks data.table::first()
```

```
## x dplyr::id()      masks plyr::id()
## x dplyr::lag()     masks stats::lag()
## x dplyr::last()    masks data.table::last()
## x dplyr::mutate()  masks plyr::mutate()
## x dplyr::rename()  masks plyr::rename()
## x dplyr::summarise() masks plyr::summarise()
## x dplyr::summarize() masks plyr::summarize()
## x purrr::transpose() masks data.table::transpose()
```

```
library("tidyr")
library("lubridate")
```

```
##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:data.table':
##
##     hour, isoweek, mday, minute, month, quarter, second, wday, week,
##     yday, year

## The following objects are masked from 'package:base':
##
##     date, intersect, setdiff, union
```

```
library("ggcorrplot")
library("ggplot2")
library("corrplot")
```

```
## corrplot 0.84 loaded
```

```
library("moments")
library('xtable')
library('countrycode')
library('class')
library("rpart")
library("rpart.plot")
library("mlbench")
library('e1071')
```

```
##
## Attaching package: 'e1071'

## The following objects are masked from 'package:moments':
##
##     kurtosis, moment, skewness
```

```
library('rpart')
library('caret')
```

```
## Loading required package: lattice
```

```
##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
## lift

library('ranger')
library('kernlab')

##
## Attaching package: 'kernlab'

## The following object is masked from 'package:purrr':
##
## cross

## The following object is masked from 'package:ggplot2':
##
## alpha

library('ggbiplot')

## Loading required package: scales

##
## Attaching package: 'scales'

## The following object is masked from 'package:kernlab':
##
## alpha

## The following object is masked from 'package:purrr':
##
## discard

## The following object is masked from 'package:readr':
##
## col_factor

## Loading required package: grid

library('ISLR')
library('devtools')

## Loading required package: usethis
```

```
# Loading the Dataset
```

```
cf_df <- read.csv(url("http://bit.ly/CarreFourDataset"))
```

Previewing the data

```
# Previewing The First Seven records in the Dataset
```

```
head(cf_df, n=7)
```

```
##      Invoice.ID Branch Customer.type Gender      Product.line Unit.price
## 1 750-67-8428      A      Member Female      Health and beauty      74.69
## 2 226-31-3081      C      Normal Female Electronic accessories      15.28
## 3 631-41-3108      A      Normal  Male      Home and lifestyle      46.33
## 4 123-19-1176      A      Member  Male      Health and beauty      58.22
## 5 373-73-7910      A      Normal  Male      Sports and travel      86.31
## 6 699-14-3026      C      Normal  Male Electronic accessories      85.39
## 7 355-53-5943      A      Member Female Electronic accessories      68.84
##      Quantity      Tax      Date Time      Payment      cogs gross.margin.percentage
## 1          7 26.1415 1/5/2019 13:08      Ewallet 522.83          4.761905
## 2          5  3.8200 3/8/2019 10:29      Cash 76.40          4.761905
## 3          7 16.2155 3/3/2019 13:23 Credit card 324.31          4.761905
## 4          8 23.2880 1/27/2019 20:33      Ewallet 465.76          4.761905
## 5          7 30.2085 2/8/2019 10:37      Ewallet 604.17          4.761905
## 6          7 29.8865 3/25/2019 18:30      Ewallet 597.73          4.761905
## 7          6 20.6520 2/25/2019 14:36      Ewallet 413.04          4.761905
##      gross.income Rating      Total
## 1      26.1415      9.1 548.9715
## 2       3.8200      9.6  80.2200
## 3      16.2155      7.4 340.5255
## 4      23.2880      8.4 489.0480
## 5      30.2085      5.3 634.3785
## 6      29.8865      4.1 627.6165
## 7      20.6520      5.8 433.6920
```

```
# Previewing The Last Seven records in the Dataset
```

```
tail(cf_df, n=7)
```

```
##      Invoice.ID Branch Customer.type Gender      Product.line Unit.price
## 994 690-01-6631      B      Normal  Male      Fashion accessories      17.49
## 995 652-49-6720      C      Member Female Electronic accessories      60.95
## 996 233-67-5758      C      Normal  Male      Health and beauty      40.35
## 997 303-96-2227      B      Normal Female      Home and lifestyle      97.38
## 998 727-02-1313      A      Member  Male      Food and beverages      31.84
## 999 347-56-2442      A      Normal  Male      Home and lifestyle      65.82
## 1000 849-09-3807      A      Member Female      Fashion accessories      88.34
##      Quantity      Tax      Date Time Payment      cogs gross.margin.percentage
## 994          10  8.7450 2/22/2019 18:35 Ewallet 174.90          4.761905
## 995           1  3.0475 2/18/2019 11:40 Ewallet  60.95          4.761905
```

```
## 996      1  2.0175 1/29/2019 13:46 Ewallet  40.35      4.761905
## 997     10 48.6900  3/2/2019 17:16 Ewallet 973.80      4.761905
## 998      1  1.5920  2/9/2019 13:22   Cash  31.84      4.761905
## 999      1  3.2910  2/22/2019 15:33   Cash  65.82      4.761905
## 1000     7 30.9190  2/18/2019 13:28   Cash 618.38      4.761905
##      gross.income Rating      Total
## 994      8.7450      6.6  183.6450
## 995      3.0475      5.9   63.9975
## 996      2.0175      6.2   42.3675
## 997     48.6900      4.4 1022.4900
## 998      1.5920      7.7   33.4320
## 999      3.2910      4.1   69.1110
## 1000     30.9190      6.6  649.2990
```

```
# Checking the Data Dimensions
```

```
dim(cf_df)
```

```
## [1] 1000  16
```

The dataset has 1000 records and 10 columns

```
# Checking the Structure of the Dataset
```

```
str(cf_df)
```

```
## 'data.frame':  1000 obs. of  16 variables:
## $ Invoice.ID      : chr  "750-67-8428" "226-31-3081" "631-41-3108" "123-19-1176" ...
## $ Branch         : chr  "A" "C" "A" "A" ...
## $ Customer.type   : chr  "Member" "Normal" "Normal" "Member" ...
## $ Gender          : chr  "Female" "Female" "Male" "Male" ...
## $ Product.line    : chr  "Health and beauty" "Electronic accessories" "Home and lifestyle" ...
## $ Unit.price      : num  74.7 15.3 46.3 58.2 86.3 ...
## $ Quantity        : int   7 5 7 8 7 7 6 10 2 3 ...
## $ Tax             : num  26.14 3.82 16.22 23.29 30.21 ...
## $ Date            : chr  "1/5/2019" "3/8/2019" "3/3/2019" "1/27/2019" ...
## $ Time            : chr  "13:08" "10:29" "13:23" "20:33" ...
## $ Payment         : chr  "Ewallet" "Cash" "Credit card" "Ewallet" ...
## $ cogs            : num  522.8 76.4 324.3 465.8 604.2 ...
## $ gross.margin.percentage: num  4.76 4.76 4.76 4.76 4.76 ...
## $ gross.income    : num  26.14 3.82 16.22 23.29 30.21 ...
## $ Rating          : num  9.1 9.6 7.4 8.4 5.3 4.1 5.8 8 7.2 5.9 ...
## $ Total           : num  549 80.2 340.5 489 634.4 ...
```

```
# Checking The Data present in each column
```

```
glimpse(cf_df)
```

```
## Rows: 1,000
## Columns: 16
```

```
## $ Invoice.ID          <chr> "750-67-8428", "226-31-3081", "631-41-3108"...
## $ Branch             <chr> "A", "C", "A", "A", "A", "C", "A", "C", "A"...
## $ Customer.type      <chr> "Member", "Normal", "Normal", "Member", "No...
## $ Gender             <chr> "Female", "Female", "Male", "Male", "Male",...
## $ Product.line       <chr> "Health and beauty", "Electronic accessorie...
## $ Unit.price         <dbl> 74.69, 15.28, 46.33, 58.22, 86.31, 85.39, 6...
## $ Quantity          <int> 7, 5, 7, 8, 7, 7, 6, 10, 2, 3, 4, 4, 5, 10,...
## $ Tax               <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ Date              <chr> "1/5/2019", "3/8/2019", "3/3/2019", "1/27/2...
## $ Time              <chr> "13:08", "10:29", "13:23", "20:33", "10:37"...
## $ Payment           <chr> "Ewallet", "Cash", "Credit card", "Ewallet"...
## $ cogs              <dbl> 522.83, 76.40, 324.31, 465.76, 604.17, 597....
## $ gross.margin.percentage <dbl> 4.761905, 4.761905, 4.761905, 4.761905, 4.7...
## $ gross.income       <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ Rating            <dbl> 9.1, 9.6, 7.4, 8.4, 5.3, 4.1, 5.8, 8.0, 7.2...
## $ Total             <dbl> 548.9715, 80.2200, 340.5255, 489.0480, 634....
```

7. Data Preparation

Uniformity

```
# Check column names
```

```
colnames(cf_df)
```

```
## [1] "Invoice.ID"          "Branch"
## [3] "Customer.type"       "Gender"
## [5] "Product.line"        "Unit.price"
## [7] "Quantity"            "Tax"
## [9] "Date"                "Time"
## [11] "Payment"             "cogs"
## [13] "gross.margin.percentage" "gross.income"
## [15] "Rating"              "Total"
```

```
# Renaming column names
```

```
names(cf_df)[1]<- 'Invoice_ID'
names(cf_df)[3]<- 'Customer_type'
names(cf_df)[5]<- 'Product_line'
names(cf_df)[6]<- 'Unit_price'
names(cf_df)[14]<- 'Gross_income'
```

```
# Checking whether the column names have been changed
```

```
colnames(cf_df)
```

We'll rename the column names for Uniformity purposes

```
## [1] "Invoice_ID"          "Branch"
## [3] "Customer_type"        "Gender"
## [5] "Product_line"         "Unit_price"
## [7] "Quantity"             "Tax"
## [9] "Date"                 "Time"
## [11] "Payment"              "cogs"
## [13] "gross.margin.percentage" "Gross_income"
## [15] "Rating"               "Total"
```

```
# Checking for the length of unique values in each column
```

```
lapply(cf_df, function (x) {length(unique(x))})
```

```
## $Invoice_ID
## [1] 1000
##
## $Branch
## [1] 3
##
## $Customer_type
## [1] 2
##
## $Gender
## [1] 2
##
## $Product_line
## [1] 6
##
## $Unit_price
## [1] 943
##
## $Quantity
## [1] 10
##
## $Tax
## [1] 990
##
## $Date
## [1] 89
##
## $Time
## [1] 506
##
## $Payment
## [1] 3
##
## $cogs
## [1] 990
##
## $gross.margin.percentage
## [1] 1
##
## $Gross_income
## [1] 990
```



```
##
## $Rating
## [1] 61
##
## $Total
## [1] 990
```

```
# Cheking if Tax and gross columns are duplicated
```

```
unique(cf_df$Tax == cf_df$Gross_income)
```

```
## [1] TRUE
```

```
# Drop the Gross Margin percentage and Tax (Tax and Gross income are duplicated ) column
```

```
cf_df <- cf_df[, -8]
cf_df <- cf_df[, -12]
```

```
dim(cf_df)
```

Gross income percentage' has one unique variable making it redundant in our analysis.

```
## [1] 1000 14
```

Completeness

```
# Checking for missing values
```

```
colSums(is.na(cf_df))
```

```
## Invoice_ID Branch Customer_type Gender Product_line
## 0 0 0 0 0
## Unit_price Quantity Date Time Payment
## 0 0 0 0 0
## cogs Gross_income Rating Total
## 0 0 0 0
```

```
# Checking for duplicate values
```




```
duplicates <- cf_df[duplicated(cf_df),]
duplicates
```

```
## [1] Invoice_ID Branch Customer_type Gender Product_line
## [6] Unit_price Quantity Date Time Payment
## [11] cogs Gross_income Rating Total
## <0 rows> (or 0-length row.names)
```

Outlier Detection

```
# Plotting boxplots for all the numerical variables
```

```
par(mfrow=c(3,2))
boxplot((cf_df$`Unit_price`), horizontal = TRUE, col = 'red', main = "Unit_price")
boxplot((cf_df$`Quantity`), horizontal = TRUE, col = 'blue', main = "Quantity")
boxplot((cf_df$`Gross`), horizontal = TRUE, col = 'green', main = "Gross_income")
boxplot((cf_df$`cogs`), horizontal = TRUE, col = 'orange', main = "cogs")
boxplot((cf_df$`Total`), horizontal = TRUE, col = 'purple', main = "Total")
boxplot((cf_df$`Rating`), horizontal = TRUE, col = 'skyblue', main = "Rating")
```



```
# Checking for anomalies in our numerical variables
#### We will not remove the outliers as they may have vital information
```

8. Exploratory Data Analysis

Univariate Analysis

```
# Checking the statistical summary of the data
```

```
summary(cf_df)
```

```
## Invoice_ID      Branch      Customer_type      Gender
## Length:1000    Length:1000  Length:1000    Length:1000
## Class :character Class :character Class :character Class :character
```

```
## Mode :character Mode :character Mode :character Mode :character
##
##
##
## Product_line      Unit_price      Quantity      Date
## Length:1000      Min. :10.08      Min. : 1.00      Length:1000
## Class :character  1st Qu.:32.88      1st Qu.: 3.00      Class :character
## Mode :character   Median :55.23      Median : 5.00      Mode :character
##                   Mean :55.67      Mean : 5.51
##                   3rd Qu.:77.94      3rd Qu.: 8.00
##                   Max. :99.96      Max. :10.00
##      Time          Payment          cogs          Gross_income
## Length:1000      Length:1000      Min. : 10.17      Min. : 0.5085
## Class :character  Class :character  1st Qu.:118.50      1st Qu.: 5.9249
## Mode :character  Mode :character  Median :241.76      Median :12.0880
##                                     Mean :307.59      Mean :15.3794
##                                     3rd Qu.:448.90      3rd Qu.:22.4453
##                                     Max. :993.00      Max. :49.6500
##      Rating          Total
## Min. : 4.000      Min. : 10.68
## 1st Qu.: 5.500      1st Qu.: 124.42
## Median : 7.000      Median : 253.85
## Mean : 6.973      Mean : 322.97
## 3rd Qu.: 8.500      3rd Qu.: 471.35
## Max. :10.000      Max. :1042.65
```

Measures of Central Tendency and Dispersion - Summary

Central Tendency - Mode, Mean and Median

```
# First, a function for mode will be created since R does not have a built in function.
```

```
getmode <- function(v) {
  uniqv <- unique(v)
  uniqv[which.max(tabulate(match(v, uniqv)))]
}
```

```
# Unit Price
```

```
mode.up <- getmode(cf_df$Unit_price)
mode.up
```

```
## [1] 83.77
```

```
mean(cf_df$Unit_price)
```

```
## [1] 55.67213
```

```
median(cf_df$Unit_price)
```

```
## [1] 55.23
```

```
# Gross_income
```

```
mode.gi <- getmode(cf_df$Gross_income)  
mode.gi
```

```
## [1] 39.48
```

```
mean(cf_df$Gross_income)
```

```
## [1] 15.37937
```

```
median(cf_df$Gross_income)
```

```
## [1] 12.088
```

```
# Quantity
```

```
mode.quan <- getmode(cf_df$Quantity)  
mode.quan
```

```
## [1] 10
```

```
mean(cf_df$Quantity)
```

```
## [1] 5.51
```

```
median(cf_df$Quantity)
```

```
## [1] 5
```

```
# Cogs
```

```
mode.cogs <- getmode(cf_df$cogs)  
mode.cogs
```

```
## [1] 789.6
```

```
mean(cf_df$cogs)
```

```
## [1] 307.5874
```

```
median(cf_df$cogs)
```

```
## [1] 241.76
```

```
# Total
mode.total <- getmode(cf_df$Total)
mode.total
```

```
## [1] 829.08
```

```
mean(cf_df$Total)
```

```
## [1] 322.9667
```

```
median(cf_df$Total)
```

```
## [1] 253.848
```

```
# Rating
mode.rating <- getmode(cf_df$Rating)
mode.rating
```

```
## [1] 6
```

```
mean(cf_df$Rating)
```

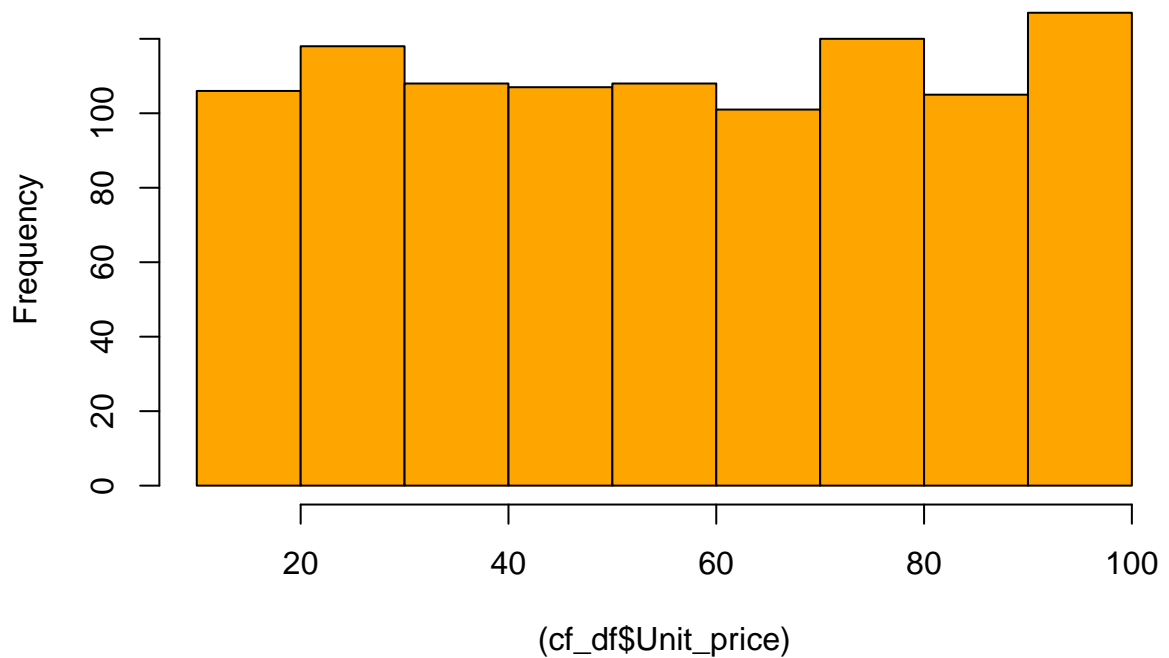
```
## [1] 6.9727
```

```
median(cf_df$Rating)
```

```
## [1] 7
```

```
# Unit price
hist((cf_df$Unit_price), col = 'orange', main = "Unit Price")
```

Measure of Dispersion and Histograms - Standard Deviation, Variance, Skewness, Kurtosis and Unit Price



Range

```
sd.up <- sd(cf_df$Unit_price)
sd.up
```

```
## [1] 26.49463
```

```
var.up <- var(cf_df$Unit_price)
var.up
```

```
## [1] 701.9653
```

```
range.up <- range(cf_df$Unit_price)
range.up
```

```
## [1] 10.08 99.96
```

```
skew.up <- skewness(cf_df$Unit_price)
skew.up
```

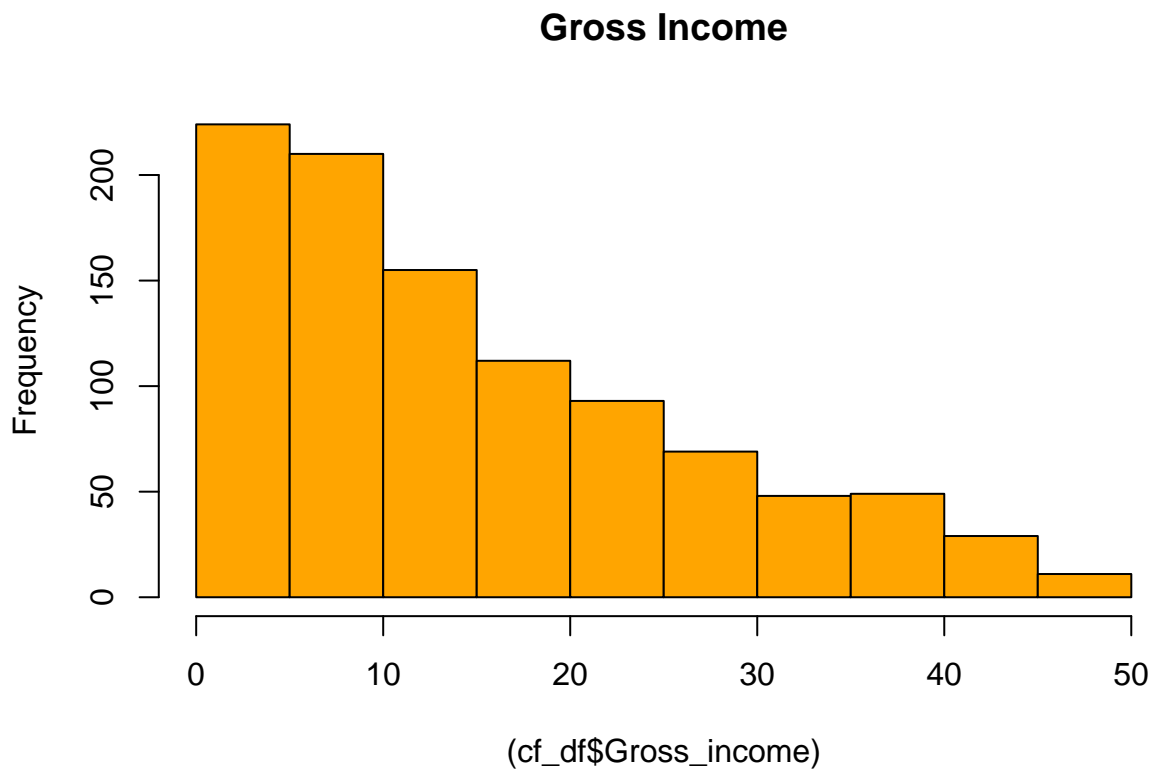
```
## [1] 0.00705623
```

```
kurt.up <- kurtosis(cf_df$Unit_price)
kurt.up
```

```
## [1] -1.222062
```

```
# Gross income
```

```
hist((cf_df$Gross_income'), col = 'orange', main = "Gross Income")
```



```
sd.gi <- sd(cf_df$Gross_income)
sd.gi
```

```
## [1] 11.70883
```

```
var.gi <- var(cf_df$Gross_income)
var.gi
```

```
## [1] 137.0966
```

```
range.gi <- range(cf_df$Gross_income)
range.gi
```

```
## [1] 0.5085 49.6500
```

```
skew.gi <- skewness(cf_df$Gross_income)
skew.gi
```

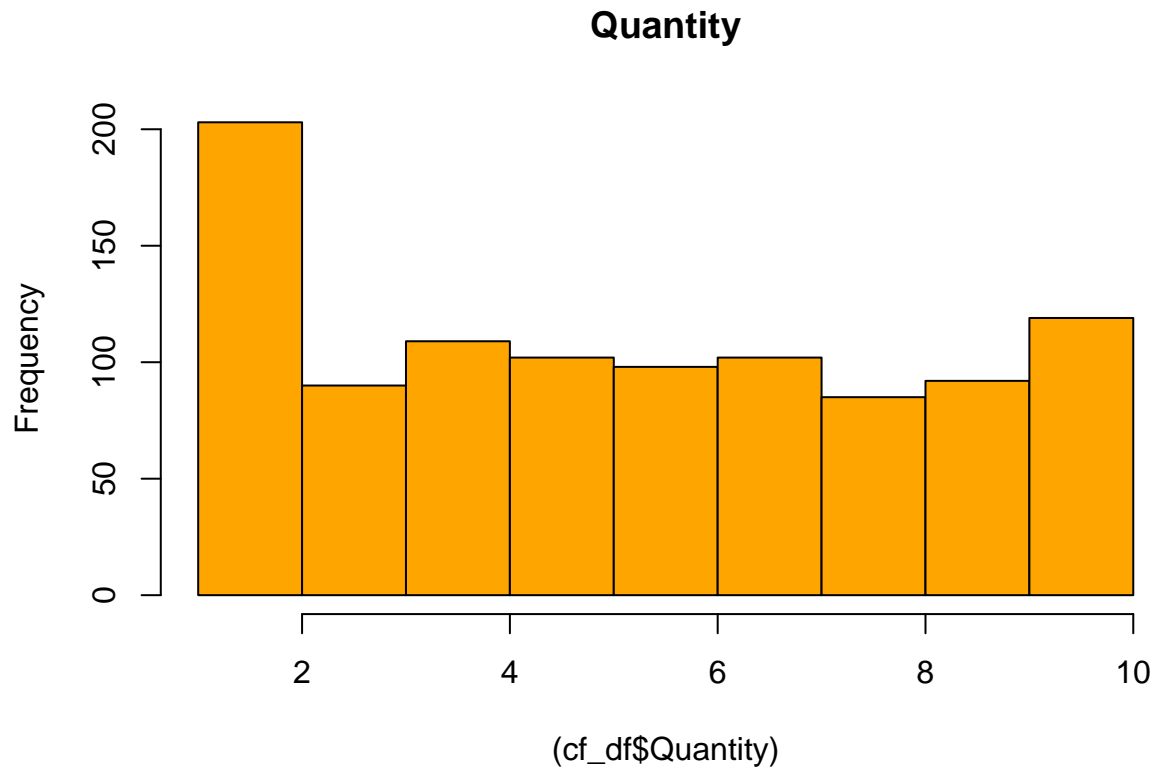
```
## [1] 0.8898939
```

```
kurt.gi <- kurtosis(cf_df$Gross_income)
kurt.gi
```

```
## [1] -0.09329206
```

```
# Quantity
```

```
hist((cf_df$`Quantity`), col = 'orange', main = "Quantity")
```



```
sd.quan <- sd(cf_df$Quantity)
sd.quan
```

```
## [1] 2.923431
```

```
var.quan <- var(cf_df$Quantity)
var.quan
```

```
## [1] 8.546446
```

```
range.quan <- range(cf_df$Quantity)
range.quan
```

```
## [1] 1 10
```




```
skew.quan <- skewness(cf_df$Quantity)
skew.quan
```

```
## [1] 0.01290225
```

```
kurt.quan <- kurtosis(cf_df$Quantity)
kurt.quan
```

```
## [1] -1.219039
```

```
# Cogs
hist((cf_df$cogs), col = 'orange', main = "Cogs")
```



```
sd.cogs <- sd(cf_df$cogs)
sd.cogs
```

```
## [1] 234.1765
```

```
var.cogs <- var(cf_df$cogs)
var.cogs
```

```
## [1] 54838.64
```

```
range.cogs <- range(cf_df$cogs)
range.cogs
```

```
## [1] 10.17 993.00
```

```
skew.cogs <- skewness(cf_df$cogs)
skew.cogs
```


```
## [1] 0.8898939
```

```
kurt.cogs <- kurtosis(cf_df$cogs)
kurt.cogs
```

```
## [1] -0.09329206
```

```
# Total
```

```
hist((cf_df$Total), col = 'orange', main = "Total")
```



```
sd.total <- sd(cf_df$Total)
sd.total
```

```
## [1] 245.8853
```

```
var.total <- var(cf_df$Total)
var.total
```

```
## [1] 60459.6
```

```
range.total <- range(cf_df$Total)
range.total
```

```
## [1] 10.6785 1042.6500
```

```
skew.total <- skewness(cf_df$Total)
skew.total
```

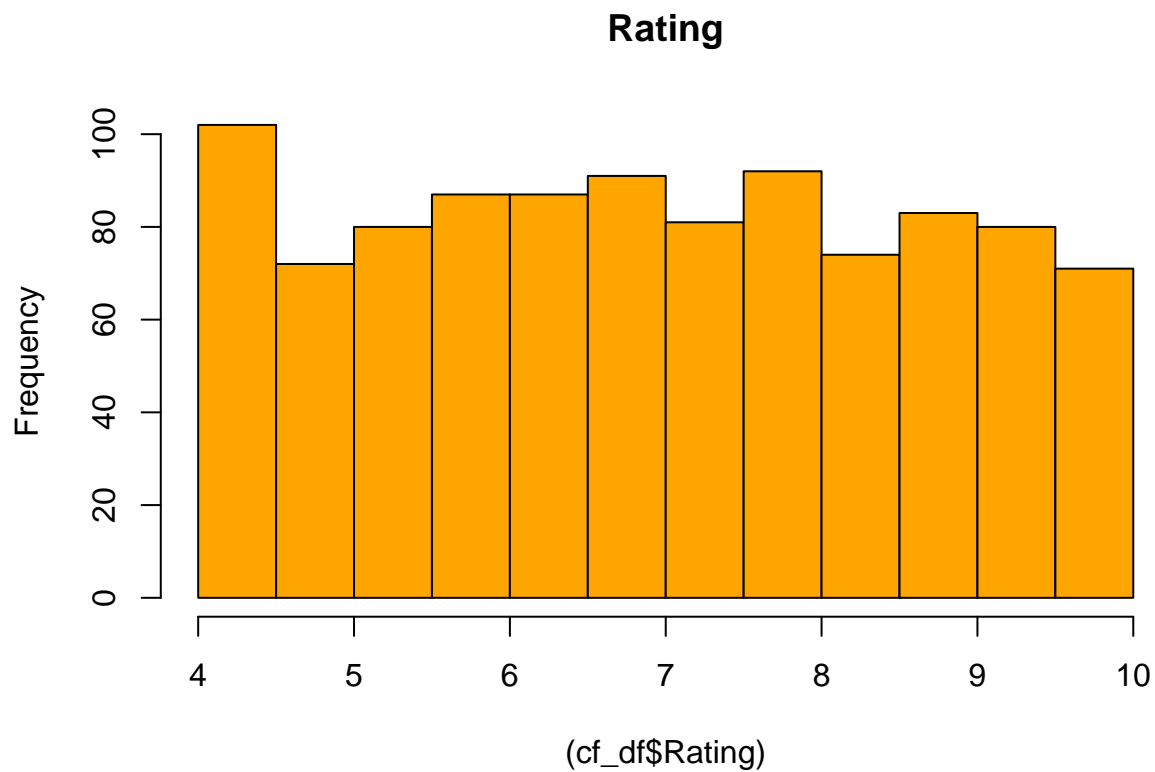
```
## [1] 0.8898939
```

```
kurt.total <- kurtosis(cf_df$Total)
kurt.total
```

```
## [1] -0.09329206
```

```
# Rating
```

```
hist((cf_df$Rating), col = 'orange', main = "Rating")
```



```
sd.r <- sd(cf_df$Rating)
sd.r
```

```
## [1] 1.71858
```

```
var.r <- var(cf_df$Rating)
var.r
```

```
## [1] 2.953518
```

```
range.r <- range(cf_df$Rating)
range.r
```

```
## [1] 4 10
```

```
skew.r <- skewness(cf_df$Rating)
skew.r
```


```
## [1] 0.008982638
```

```
kurt.r <- kurtosis(cf_df$Rating)
kurt.r
```

```
## [1] -1.155525
```


Barplots

```
# Branches
barplot(table(cf_df$Branch), main = "Branches")
```

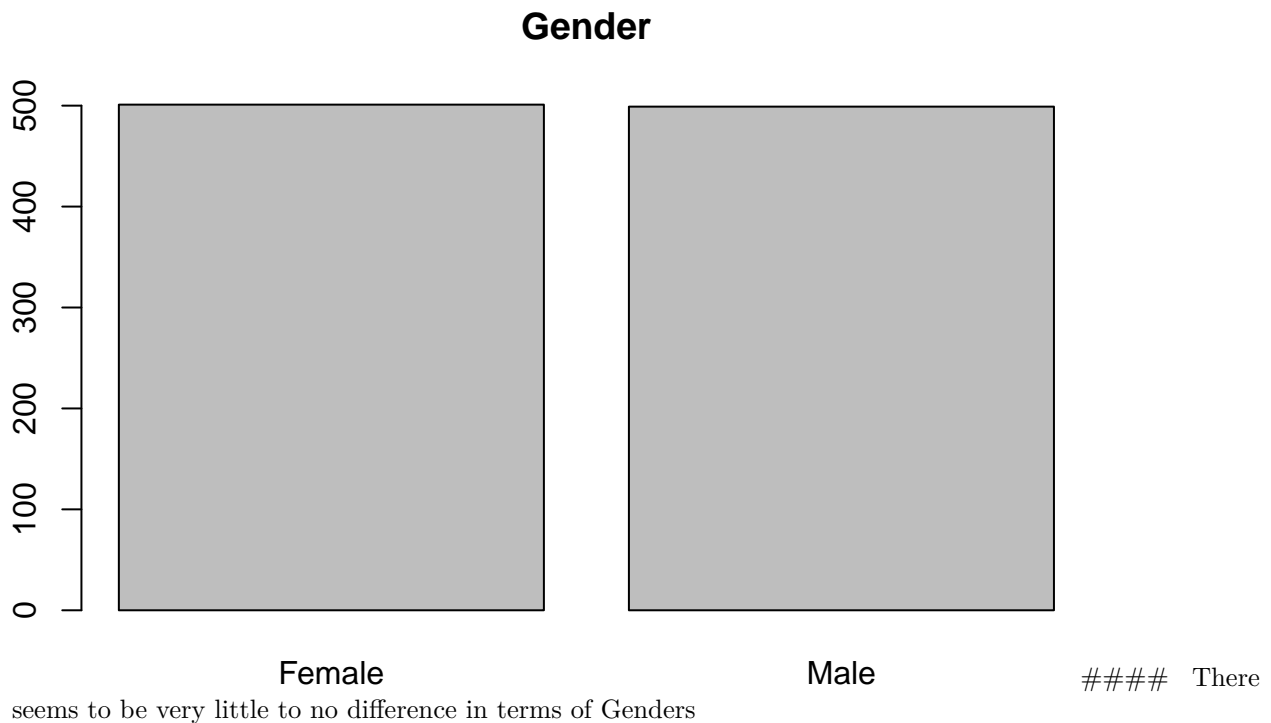


A has the most activity but there is very little difference between them.

```
# Customer Type
barplot(table(cf_df$Customer_type), main = "Customer Type")
```

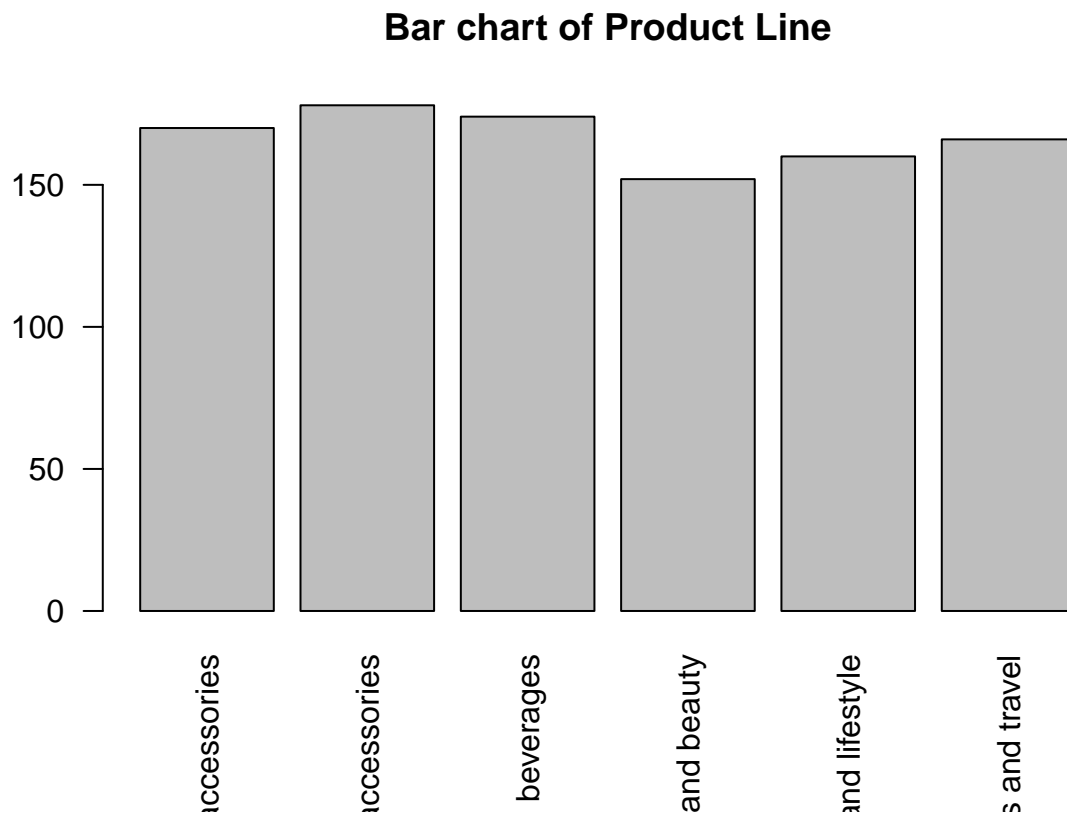


```
# Gender
barplot(table(cf_df$Gender), main = "Gender")
```



```
# Product Line
```

```
barplot(table(cf_df$Product_line), main = "Bar chart of Product Line", las=2)
```




It seems

Electronics and accessories and food and beverages are the most popular product lines.

```
# Payment
```

```
barplot(table(cf_df$Payment), main = "Payment")
```



Bivariate Analysis

```
# Creating a new dataframe num with numerical data variables
```

```
Unit_price<- cf_df$Unit_price
Gross_income<-cf_df$Gross_income
cogs<-cf_df$cogs
Total<-cf_df$Total
Rating<-cf_df$Rating
```

```
num_data <- data.frame(Unit_price, Gross_income, cogs, Total, Rating)
head(num_data)
```

```
##   Unit_price Gross_income   cogs   Total Rating
## 1    74.69      26.1415 522.83 548.9715    9.1
## 2    15.28       3.8200  76.40  80.2200    9.6
## 3    46.33      16.2155 324.31 340.5255    7.4
## 4    58.22      23.2880 465.76 489.0480    8.4
## 5    86.31      30.2085 604.17 634.3785    5.3
## 6    85.39      29.8865 597.73 627.6165    4.1
```

```
# Correlation
```

```
# Correlation is a statistical technique that can show whether and how strongly pairs of variables are
```

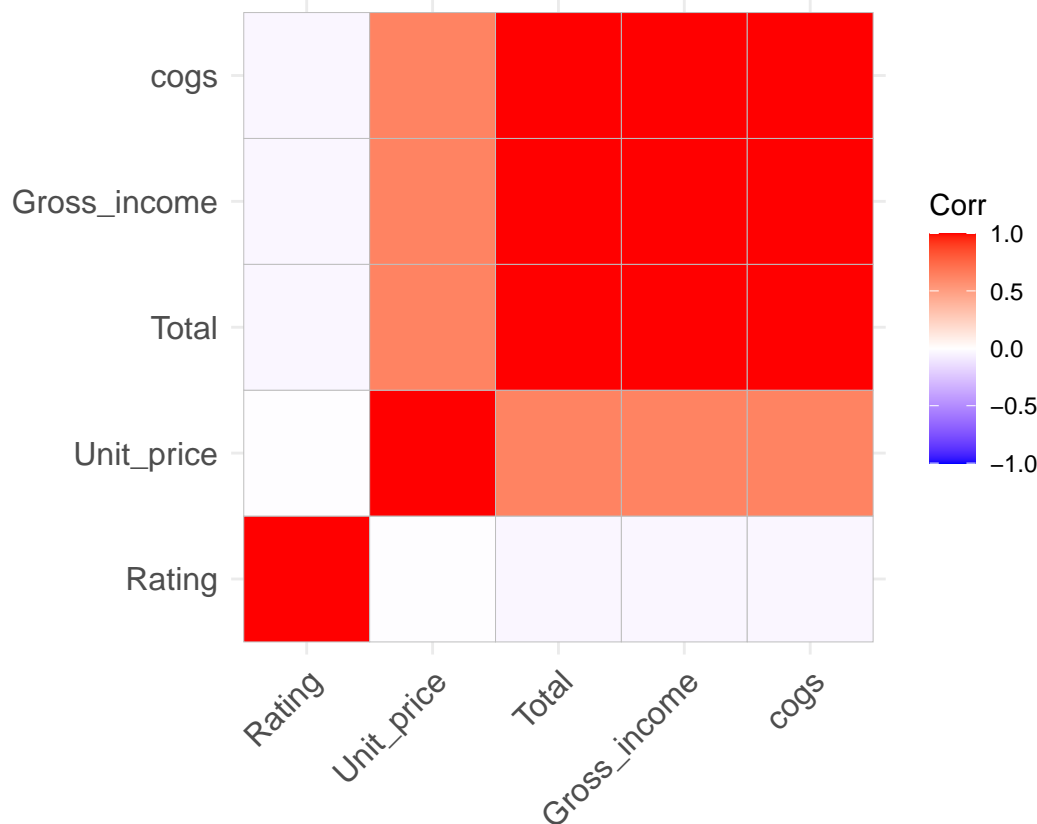
```
# Calculating the correlation matrix
```

```
corr <- cor(num_data)
head(corr)
```

```
##          Unit_price Gross_income      cogs      Total      Rating
## Unit_price  1.000000000  0.6339621  0.6339621  0.6339621 -0.008777507
## Gross_income 0.633962089  1.0000000  1.0000000  1.0000000 -0.036441705
## cogs         0.633962089  1.0000000  1.0000000  1.0000000 -0.036441705
## Total        0.633962089  1.0000000  1.0000000  1.0000000 -0.036441705
## Rating       -0.008777507 -0.0364417 -0.0364417 -0.0364417  1.000000000
```

```
# Plotting the correlation matrix
```

```
ggcorrplot(corr,hc.order = TRUE)
```




We observe that most of the variables are perfectly correlated which is problematic in modelling hence the need for feature extraction or feature selection.


Scatterplots

```
par(mfrow=c(2,4))
plot(num_data$Gross_income,num_data$Total, main="Total vs. Gross income")
plot(num_data$cogs, num_data$Total, main="Total vs. cogs")
plot(num_data$Unit_price, num_data$Total, main="Total vs. Unit price")
plot(num_data$Unit_price,num_data$cogs, main="Unit price vs. cogs")
plot(num_data$Unit_price,num_data$Gross_income, main="Unit price vs. Gross_income")
plot(num_data$Unit_price,num_data$Total, main="Unit price vs. Total")
```



Total vs. Gross income




Total vs. cogs




Total vs. Unit price




Unit price vs. cogs



Unit price vs. Gross_income



Unit price vs. Total



9. Implementing The Solution

Principal Component Analysis

We'll perform and visualize PCA in the given dataset.

```
# Selecting the numerical data
```

```
cf_df_num <- select_if(cf_df, is.numeric)
str(cf_df_num)
```

```
## 'data.frame': 1000 obs. of 6 variables:
## $ Unit_price : num 74.7 15.3 46.3 58.2 86.3 ...
## $ Quantity : int 7 5 7 8 7 7 6 10 2 3 ...
## $ cogs : num 522.8 76.4 324.3 465.8 604.2 ...
## $ Gross_income: num 26.14 3.82 16.22 23.29 30.21 ...
## $ Rating : num 9.1 9.6 7.4 8.4 5.3 4.1 5.8 8 7.2 5.9 ...
## $ Total : num 549 80.2 340.5 489 634.4 ...
```

```
# We then pass the data to the prcomp() and set the center and scale arguments, to be FALSE and TRUE
```

```
ef.pca <- prcomp(cf_df_num, center = FALSE, scale. = TRUE)
summary(ef.pca)
```

```
## Importance of components:
```


```
##          PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation 2.3249 0.60849 0.44308 0.16806 3.184e-16 1.283e-16
## Proportion of Variance 0.9009 0.06171 0.03272 0.00471 0.000e+00 0.000e+00
## Cumulative Proportion 0.9009 0.96257 0.99529 1.00000 1.000e+00 1.000e+00
```

```
# Calling str() to have a look at our PCA object
str(ef.pca)
```

PC1 explains 90% of the total variance, which means that more than three-quarters of the information in the dataset can be encapsulated by just that one Principal Component. PC2 explains 6.1% of the variance, PC3 - 2.2% and PC4- 0.4%

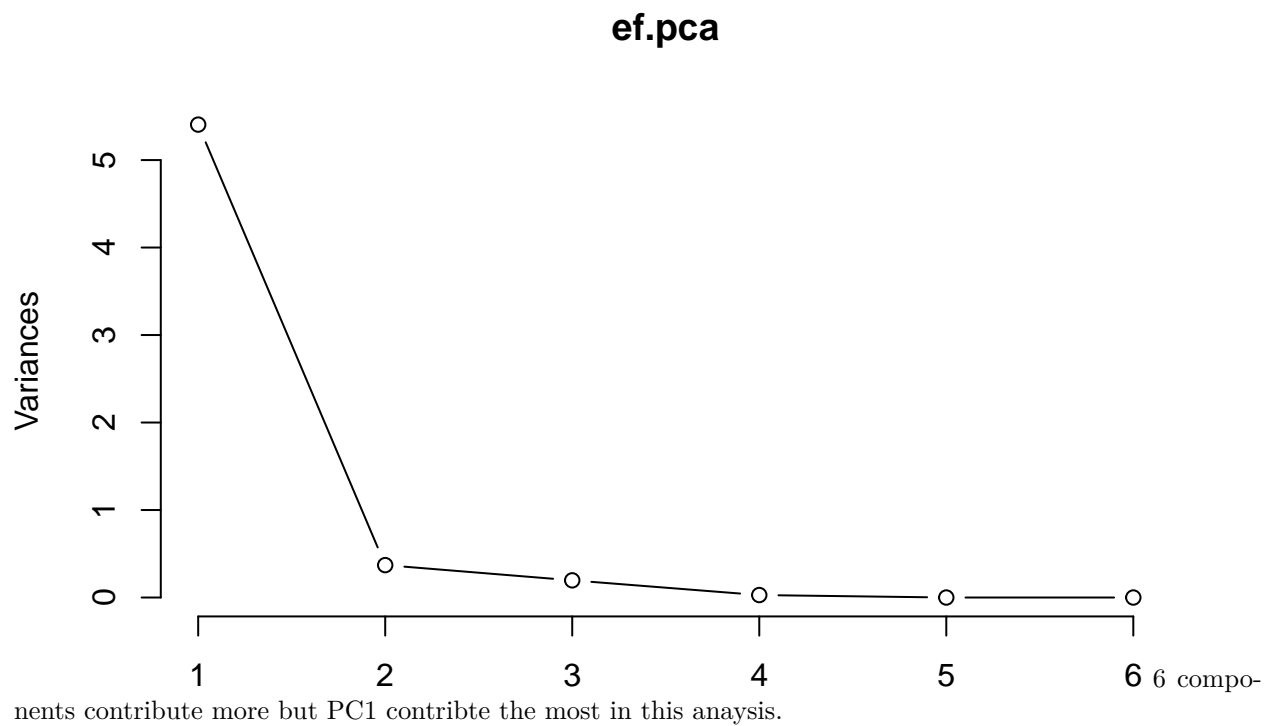
```
## List of 5
## $ sdev      : num [1:6] 2.32 6.08e-01 4.43e-01 1.68e-01 3.18e-16 ...
## $ rotation: num [1:6, 1:6] -0.402 -0.405 -0.421 -0.421 -0.379 ...
## .. attr(*, "dimnames")=List of 2
## .. ..$ : chr [1:6] "Unit_price" "Quantity" "cogs" "Gross_income" ...
## .. ..$ : chr [1:6] "PC1" "PC2" "PC3" "PC4" ...
## $ center   : logi FALSE
## $ scale    : Named num [1:6] 61.68 6.24 386.71 19.34 7.18 ...
## .. attr(*, "names")= chr [1:6] "Unit_price" "Quantity" "cogs" "Gross_income" ...
## $ x        : num [1:1000, 1:6] -3.13 -1.18 -2.2 -2.86 -3.27 ...
## .. attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:6] "PC1" "PC2" "PC3" "PC4" ...
## - attr(*, "class")= chr "prcomp"
```

```
# We will now plot our pca.
set.seed(123)
ggbiplot(ef.pca, labels=rownames(ef.pca),ellipse = TRUE,obs.scale=1,var.scale=1)
```



This is not so clear, we'll therefore plot to see the number of components that contribute more to PC1

```
plot(ef.pca, type="l")
```



Part 2: Feature Selection

```
# Calculating the correlation matrix
corrmatrix <- cor(num_data)

# Find attributes that are highly correlated
highcorr <- findCorrelation(corrmatrix, cutoff=0.75)

# Highly correlated attributes
highcorr
```

```
## [1] 2 3
```


```
names(num_data[highcorr])
```

```
## [1] "Gross_income" "cogs"
```

```
# Removing Redundant Features
num_data_clean <- num_data[-highcorr]
```

```
# Performing our graphical comparison
par(mfrow = c(1, 2))
corrplot(corrmatrix, order = "hclust")
corrplot(cor(num_data_clean), order = "hclust")
```

We observe that “Gross_income” and “cogs” are highly correlated with the other features.



Removing highly correlated variables result to less corelated variables. Hence the selected features are Unit_Price, Total and Rating. There are no more highly correlated variables.

Part 3 : Association Rules

```
# Installing and reading the necessary packages for the association rules analysis
#install.packages("arules", dependencies = TRUE)
library(arules)
```

This section will require that you create association rules that will allow you to identify relationships between variables in the dataset.

```
## Loading required package: Matrix

##
## Attaching package: 'Matrix'

## The following objects are masked from 'package:tidyr':
##
##   expand, pack, unpack

##
## Attaching package: 'arules'

## The following object is masked from 'package:kernlab':
##
##   size
```

```
## The following object is masked from 'package:dplyr':
##
##   recode
```

```
## The following objects are masked from 'package:base':
##
##   abbreviate, write
```

```
# Reading and previewing the dataset as transactions
```

```
sdf <- read.transactions("http://bit.ly/SupermarketDatasetII")
```

```
## Warning in asMethod(object): removing duplicated items in transactions
```

```
head(sdf)
```

```
## transactions in sparse format with
## 6 transactions (rows) and
## 5729 items (columns)
```

```
# Checking the dimensions of the data
dim(sdf)
```

```
## [1] 7501 5729
```

```
# Displaying the structure of our dataset
```

```
str(sdf)
```

The dataset has 7,501 transactions and 5729 columns

```
## Formal class 'transactions' [package "arules"] with 3 slots
##   ..@ data      :Formal class 'ngCMatrix' [package "Matrix"] with 5 slots
##   .. .. ..@ i    : int [1:23299] 1087 1614 1705 1732 1993 2101 2105 2358 2444 3463 ...
##   .. .. ..@ p    : int [1:7502] 0 15 16 17 18 24 27 31 33 36 ...
##   .. .. ..@ Dim   : int [1:2] 5729 7501
##   .. .. ..@ Dimnames:List of 2
##   .. .. .. ..$ : NULL
##   .. .. .. ..$ : NULL
##   .. .. ..@ factors : list()
##   ..@ itemInfo      :'data.frame': 5729 obs. of 1 variable:
##   .. ..$ labels: chr [1:5729] "&" "accessories" "accessories,antioxydant" "accessories,champagne,fre
##   ..@ itemsetInfo:'data.frame': 0 obs. of 0 variables
```

```
# Verifying the class of the object
class(sdf)
```

```
## [1] "transactions"
## attr(,"package")
## [1] "arules"
```

```
# Previewing the first few items in the dataset
inspect(sdf[1:5])
```

```
##      items
## [1] {cheese,energy,
##      drink,tomato,
##      fat,
##      flour,yams,cottage,
##      grapes,whole,
##      juice,frozen,
##      juice,low,
##      mix,green,
##      oil,
##      shrimp,almonds,avocado,vegetables,
##      smoothie,spinach,olive,
##      tea,honey,salad,mineral,
##      water,salmon,antioxydant,
##      weat,
##      yogurt,green}
## [2] {burgers,meatballs,eggs}
## [3] {chutney}
## [4] {turkey,avocado}
## [5] {bar,whole,
##      mineral,
##      rice,green,
##      tea,
##      water,milk,energy,
##      wheat}
```

```
# Previewing the items in the dataset as if it were in a dataframe
items<-as.data.frame(itemLabels(sdf))
colnames(items) <- "Item"
head(items, 10)
```

```
##      Item
## 1      &
## 2      accessories
## 3      accessories,antioxydant
## 4      accessories,champagne,fresh
## 5      accessories,champagne,protein
## 6      accessories,chocolate
## 7      accessories,chocolate,champagne,frozen
## 8      accessories,chocolate,frozen
## 9      accessories,chocolate,low
## 10     accessories,chocolate,pasta,salt
```

```
# Getting the summary statistics of the data

summary(sdf)
```

```
## transactions as itemMatrix in sparse format with
## 7501 rows (elements/itemsets/transactions) and
```



```
## 5729 columns (items) and a density of 0.0005421748
##
## most frequent items:
##      tea   wheat mineral      fat  yogurt (Other)
##      803     645     577     574     543    20157
##
## element (itemset/transaction) length distribution:
## sizes
##      1     2     3     4     5     6     7     8     9    10    11    12    13    15    16
## 1603 2007 1382  942  651  407  228  151   70   39   13    5    1    1    1
##
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
##      1.000   2.000   3.000   3.106   4.000   16.000
##
## includes extended item information - examples:
##                      labels
## 1                      &
## 2                  accessories
## 3 accessories,antioxydant
```

```
# Exploring the frequency of some articles
```

```
itemFrequency(sdf[, 7:11],type = "absolute")
```

```
## accessories,chocolate,champagne,frozen      accessories,chocolate,frozen
##                      1                      1
##      accessories,chocolate,low      accessories,chocolate,pasta,salt
##                      1                      1
##      accessories,chocolate,salt,green
##                      1
```

```
# Producing a chart of frequencies and filtering to consider only items with a minimum percentage of support
par(mfrow = c(1, 3))
# plot the frequency of items
itemFrequencyPlot(sdf, topN = 10,col="grey")
itemFrequencyPlot(sdf, support = 0.1,col="pink")
```



The top 10 most common items in the transactions dataset are Mineral water, eggs, spaghetti, french fries, chocolate, green tea, milk, ground beef, frozen vegetables and pancakes.

```
# Setting the parameters for our association analysis
```

```
rules <- apriori (sdf, parameter = list(supp = 0.001, conf = 0.8))
```

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##          0.8    0.1    1 none FALSE             TRUE     5   0.001    1
## maxlen target  ext
##      10  rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 7
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[5729 item(s), 7501 transaction(s)] done [0.01s].
## sorting and recoding items ... [354 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 done [0.00s].
## writing ... [271 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

```
rules
```

```
## set of 271 rules
```



```
# Building a apriori model with Min Support as 0.002 and confidence as 0.8.
```

```
rules2 <- apriori (sdf,parameter = list(supp = 0.002, conf = 0.8))
```

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##           0.8   0.1   1 none FALSE                TRUE     5   0.002     1
## maxlen target  ext
##       10  rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 15
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[5729 item(s), 7501 transaction(s)] done [0.01s].
## sorting and recoding items ... [189 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 done [0.00s].
## writing ... [99 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

```
rules2
```

```
## set of 99 rules
```

```
# Building apriori model with Min Support as 0.002 and confidence as 0.6.
```

```
rules3 <- apriori (sdf, parameter = list(supp = 0.001, conf = 0.6))
```

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##           0.6   0.1   1 none FALSE                TRUE     5   0.001     1
## maxlen target  ext
##       10  rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 7
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[5729 item(s), 7501 transaction(s)] done [0.02s].
## sorting and recoding items ... [354 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
```

```
## checking subsets of size 1 2 3 4 done [0.00s].
## writing ... [319 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

```
rules3
```

```
## set of 319 rules
```

```
#Performing an exploration of our model using the summary function
summary(rules)
```

```
## set of 271 rules
##
## rule length distribution (lhs + rhs):sizes
##   2   3   4
## 107 144  20
##
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   2.000   2.000   3.000   2.679   3.000   4.000
##
## summary of quality measures:
##      support      confidence      coverage      lift
##   Min.   :0.001067   Min.   :0.800   Min.   :0.001067   Min.   : 7.611
##   1st Qu.:0.001200   1st Qu.:0.931   1st Qu.:0.001200   1st Qu.: 11.630
##   Median :0.001600   Median :1.000   Median :0.001600   Median : 13.068
##   Mean   :0.002834   Mean   :0.963   Mean   :0.002973   Mean   : 22.372
##   3rd Qu.:0.002666   3rd Qu.:1.000   3rd Qu.:0.002800   3rd Qu.: 20.218
##   Max.   :0.068391   Max.   :1.000   Max.   :0.076523   Max.   :613.718
##      count
##   Min.   : 8.00
##   1st Qu.: 9.00
##   Median :12.00
##   Mean   :21.26
##   3rd Qu.:20.00
##   Max.   :513.00
##
## mining info:
##   data ntransactions support confidence
##   sdf           7501   0.001         0.8
```

```
# Observing rules built in our model i.e. first 10 model rules
```

```
inspect(rules[1:10])
```

```
##      lhs      rhs      support      confidence
## [1] {cookies,low} => {yogurt} 0.001066524 1.0
## [2] {cookies,low} => {fat} 0.001066524 1.0
## [3] {extra}      => {dark} 0.001066524 1.0
## [4] {burgers,whole} => {wheat} 0.001199840 1.0
## [5] {fries,escalope,pasta,mushroom} => {cream} 0.001066524 1.0
## [6] {fries,cookies,green} => {tea} 0.001333156 1.0
## [7] {shrimp,whole} => {wheat} 0.001066524 1.0
```

```
## [8] {rice, cake} => {wheat} 0.001333156 1.0
## [9] {tomatoes, whole} => {wheat} 0.001066524 0.8
## [10] {rice, chocolate} => {wheat} 0.001199840 0.9
##      coverage    lift      count
## [1] 0.001066524 13.813996    8
## [2] 0.001066524 13.067944    8
## [3] 0.001066524 83.344444    8
## [4] 0.001199840 11.629457    9
## [5] 0.001066524 47.777070    8
## [6] 0.001333156  9.341220   10
## [7] 0.001066524 11.629457    8
## [8] 0.001333156 11.629457   10
## [9] 0.001333156  9.303566    8
## [10] 0.001333156 10.466512    9
```

```
# Ordering these rules by a criteria such as the level of confidence then looking at the first five rules
rules<-sort(rules, by="confidence", decreasing=TRUE)
inspect(rules[1:5])
```

```
##      lhs                                rhs      support    confidence
## [1] {cookies, low}                      => {yogurt} 0.001066524 1
## [2] {cookies, low}                      => {fat}    0.001066524 1
## [3] {extra}                             => {dark}   0.001066524 1
## [4] {burgers, whole}                    => {wheat}  0.001199840 1
## [5] {fries, escalope, pasta, mushroom} => {cream}  0.001066524 1
##      coverage    lift      count
## [1] 0.001066524 13.81400    8
## [2] 0.001066524 13.06794    8
## [3] 0.001066524 83.34444    8
## [4] 0.001199840 11.62946    9
## [5] 0.001066524 47.77707    8
```

Part 4 : Anomaly Detection

```
# Reading and previewing the dataset
ano <- read.csv("http://bit.ly/CarreFourSalesDataset")
head(ano)
```

This is to detect whether there are any anomalies in the given sales dataset.

```
##      Date      Sales
## 1 1/5/2019 548.9715
## 2 3/8/2019 80.2200
## 3 3/3/2019 340.5255
## 4 1/27/2019 489.0480
## 5 2/8/2019 634.3785
## 6 3/25/2019 627.6165
```

```
# Checking the dimensions of our dataset
```

```
dim(ano)
```

```
## [1] 1000    2
```

```
# Checking the structure of our dataset
```

```
str(ano)
```

```
## 'data.frame':    1000 obs. of  2 variables:  
## $ Date : chr  "1/5/2019" "3/8/2019" "3/3/2019" "1/27/2019" ...  
## $ Sales: num  549 80.2 340.5 489 634.4 ...
```

```
# Checking the Statistical summary
```

```
summary(ano)
```

```
##      Date           Sales  
## Length:1000      Min.   : 10.68  
## Class :character  1st Qu.: 124.42  
## Mode  :character  Median : 253.85  
##                               Mean  : 322.97  
##                               3rd Qu.: 471.35  
##                               Max.   :1042.65
```

```
# Checking for null values
```

```
colSums(is.na(ano))
```

```
## Date Sales  
##    0     0
```

```
# Changing the date column from character to Date
```

```
ano <- transform(ano, Date = format(as.Date(Date, '%m/%d/%Y'), '%Y/%m/%d'))  
ano <- transform(ano, Date = as.Date(Date))  
sapply(ano, class)
```

```
##      Date      Sales  
## "Date" "numeric"
```

```
# Grouping the dataset by the Date column
```

```
Sales <- ano$Sales  
Date <- ano$Date  
ano = ano %>% arrange(Date)  
head(ano)
```

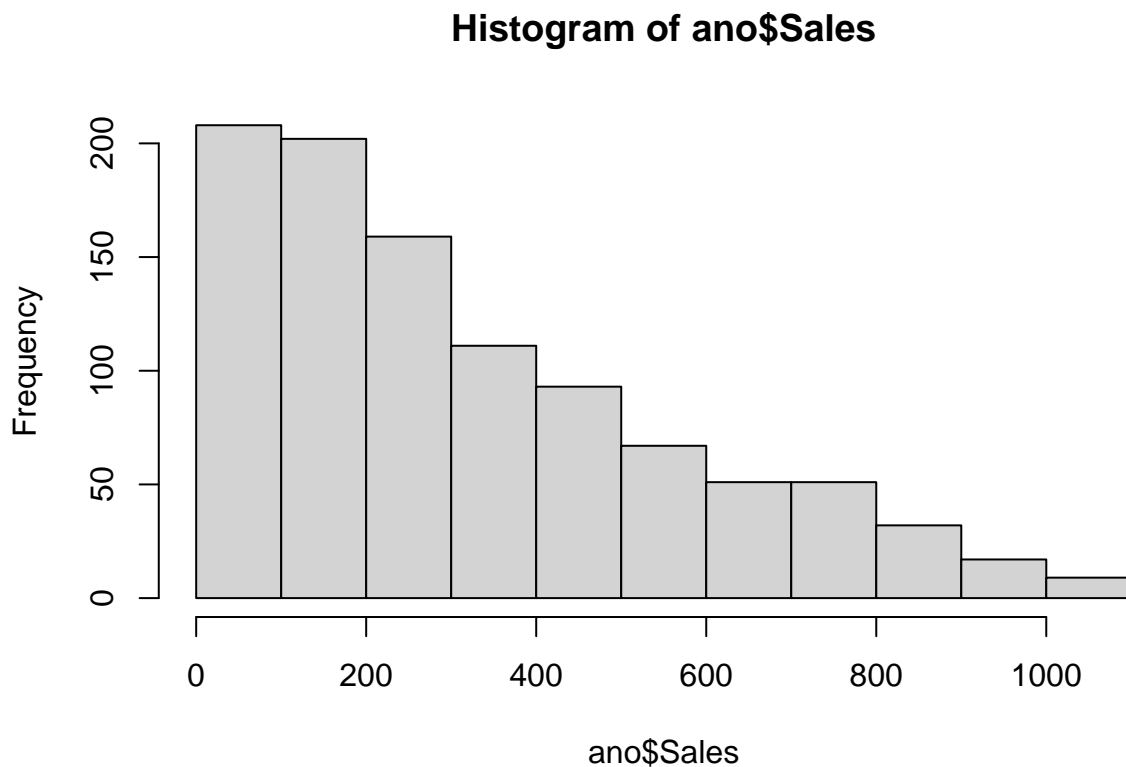
```
##      Date   Sales
## 1 2019-01-01 457.443
## 2 2019-01-01 399.756
## 3 2019-01-01 470.673
## 4 2019-01-01 388.290
## 5 2019-01-01 132.762
## 6 2019-01-01 132.027
```

```
# Sales Distribution
# sort dates in ascending order
ano_sales <- ano[order(ano$Date),]
head(ano_sales, 5)
```

```
##      Date   Sales
## 1 2019-01-01 457.443
## 2 2019-01-01 399.756
## 3 2019-01-01 470.673
## 4 2019-01-01 388.290
## 5 2019-01-01 132.762
```

```
## Plotting Histogram to show sales distribution

hist(ano$Sales)
```



```
# Transactions Count
# We'll group the transactions per day then tally them

ano_count <- ano %>% group_by(Date) %>% tally()
```

```
colnames(ano_count) <- c('Date', 'Count')
head(ano_count)
```

```
## # A tibble: 6 x 2
##   Date      Count
##   <date>    <int>
## 1 2019-01-01     12
## 2 2019-01-02      8
## 3 2019-01-03      8
## 4 2019-01-04      6
## 5 2019-01-05     12
## 6 2019-01-06      9
```

```
## Visualizing Transaction count .
library(anomalize)
```

```
## == Use anomalize to improve your Forecasts by 50%! =====
## Business Science offers a 1-hour course - Lab #18: Time Series Anomaly Detection!
## </> Learn more at: https://university.business-science.io/p/learning-labs-pro </>
```

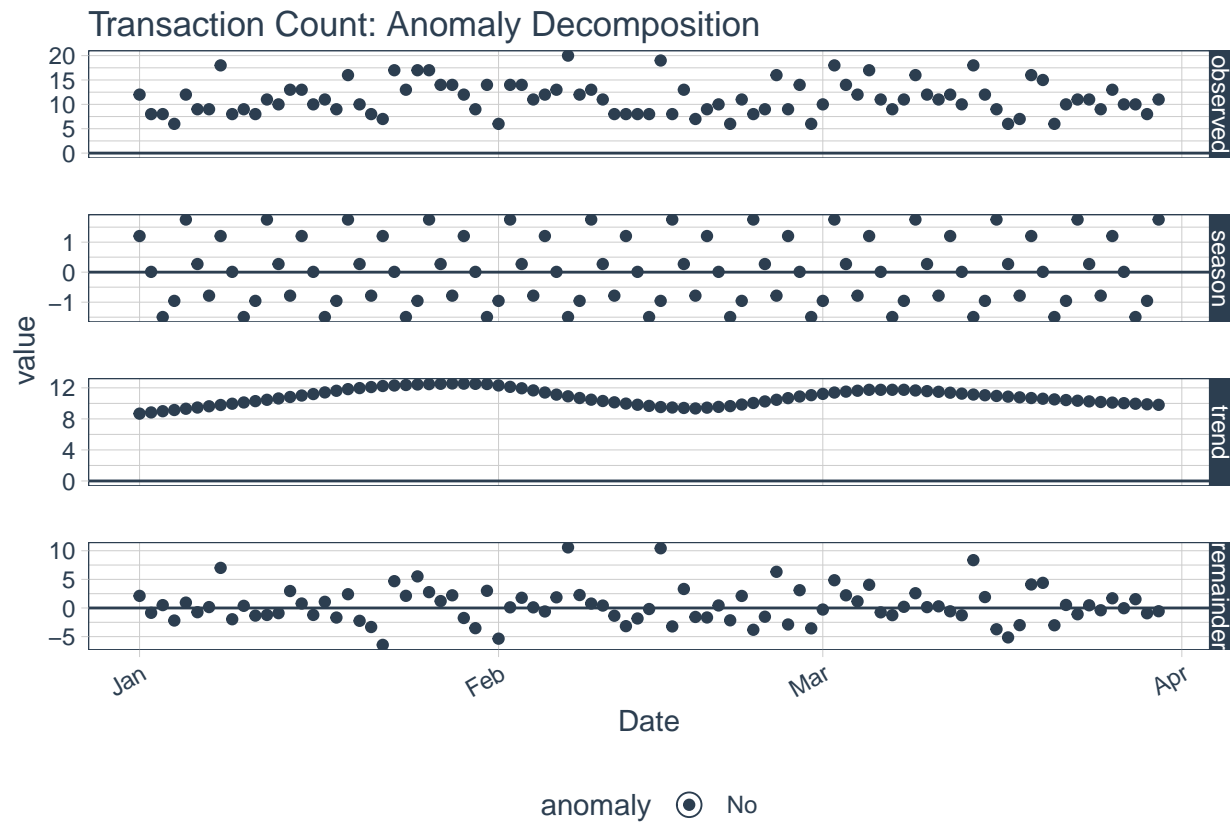
```
ano_count %>%
  time_decompose(Count) %>%
  anomalize(remainder) %>%
  plot_anomaly_decomposition() +
  ggtitle("Transaction Count: Anomaly Decomposition")
```

```
## Converting from tbl_df to tbl_time.
## Auto-index message: index = Date
```

```
## frequency = 7 days
```

```
## trend = 30 days
```

```
## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo
```




```
# Visualizing Transaction count.
ano_count %>%
  time_decompose(Count) %>%
  anomalise(remainder) %>%
  time_recompose() %>%
  plot_anomalies(time_recomposed = TRUE, ncol = 3, alpha_dots = 0.5) +
  ggtitle("Transaction Count: Anomaly Recomposition")
```

```
## Converting from tbl_df to tbl_time.
## Auto-index message: index = Date
```

```
## frequency = 7 days
```

```
## trend = 30 days
```



There were no anomalies that were detected in the number of transactions done per day between January and April.